Lab Setup Instructions

The Lab Environment

The environment you will be using for the labs is a Horton Data Platform (HDP) virtual machine running inside an onlune Nuvelab virtual machine. In the lab materials, the following terms are used consistently for clarity.

1. Host Machine. This is the machine you are accessing the online Nuvelab environment from, usually your normal work computer or laptop.
2. Nuvelab VM. This is the Nuvelab virtual machine you initially log into.
3. HPD VM. This is the Virtual Box VM that contains the HDP

Starting the Lab Environment

1. Using the credential provided, log into the Nuvelab site. You should see your Nuvelab VM in an off state.

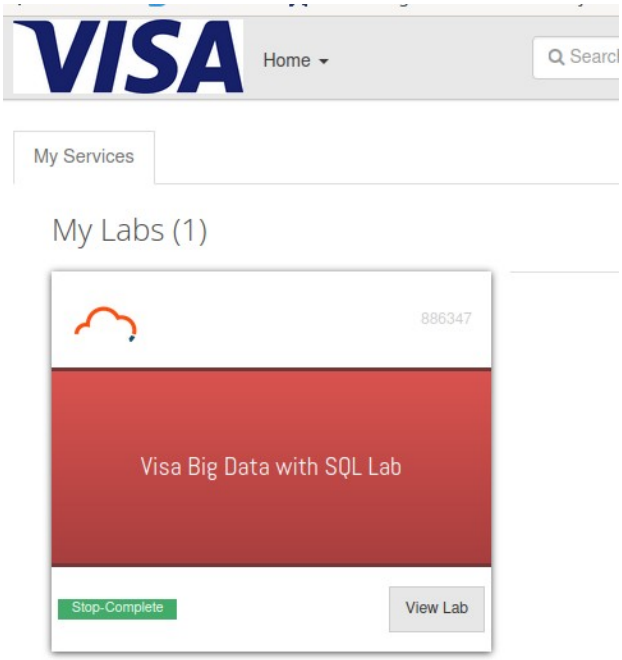

2. Click on the Lab Machine to get to the start button

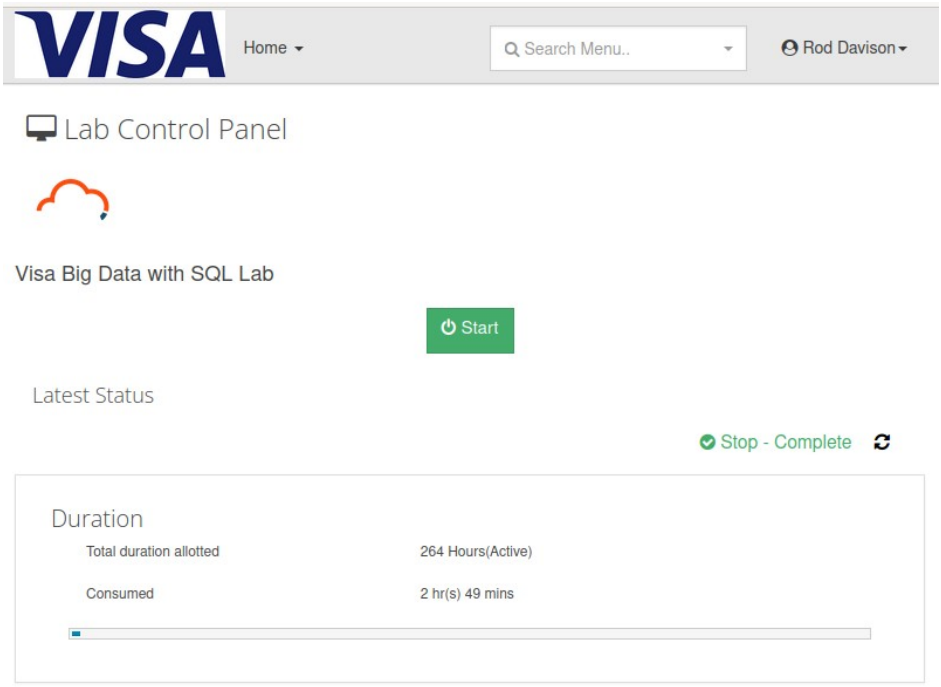

3. Press the start button. It may take a while to boot up but eventually you will see a link to the URL you will use to access the Nuvelab VM.

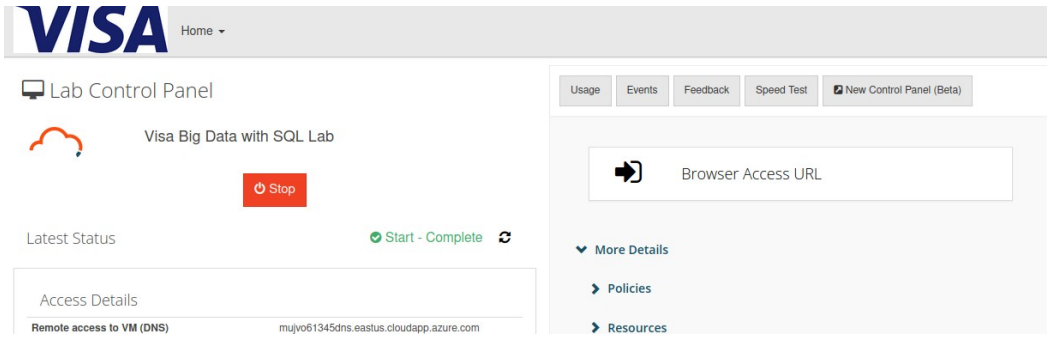

4. Going to the URL will bring up the Nuvelab VM desk top. *If you get a dialog box asking for credentials at this point, you are already logged in so just click on cancel.*
5. Double click on the “Oracle VM Virtual Box” icon to start Virtual Box. You should see the start screen below.

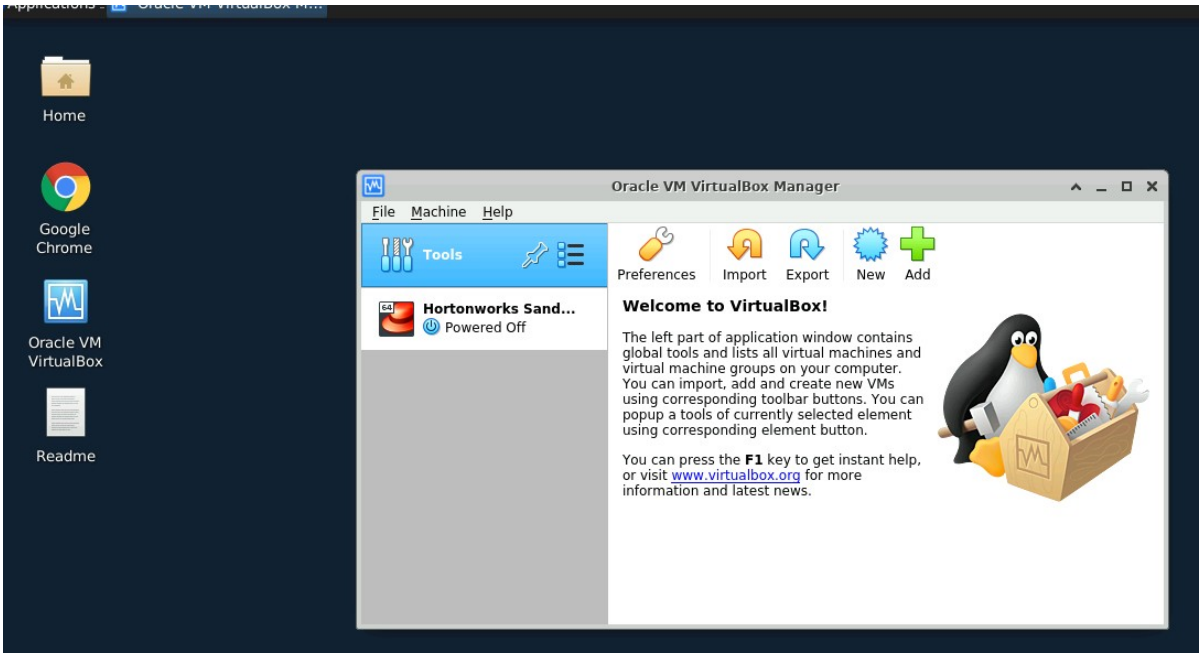

6. Select the Nortonworks Sandbox entry and a “Start” green arrow should appear which you can use to start the HPD VM.

7. The HPD VM may take a while to start up, but eventually you will see a display that looks like this.

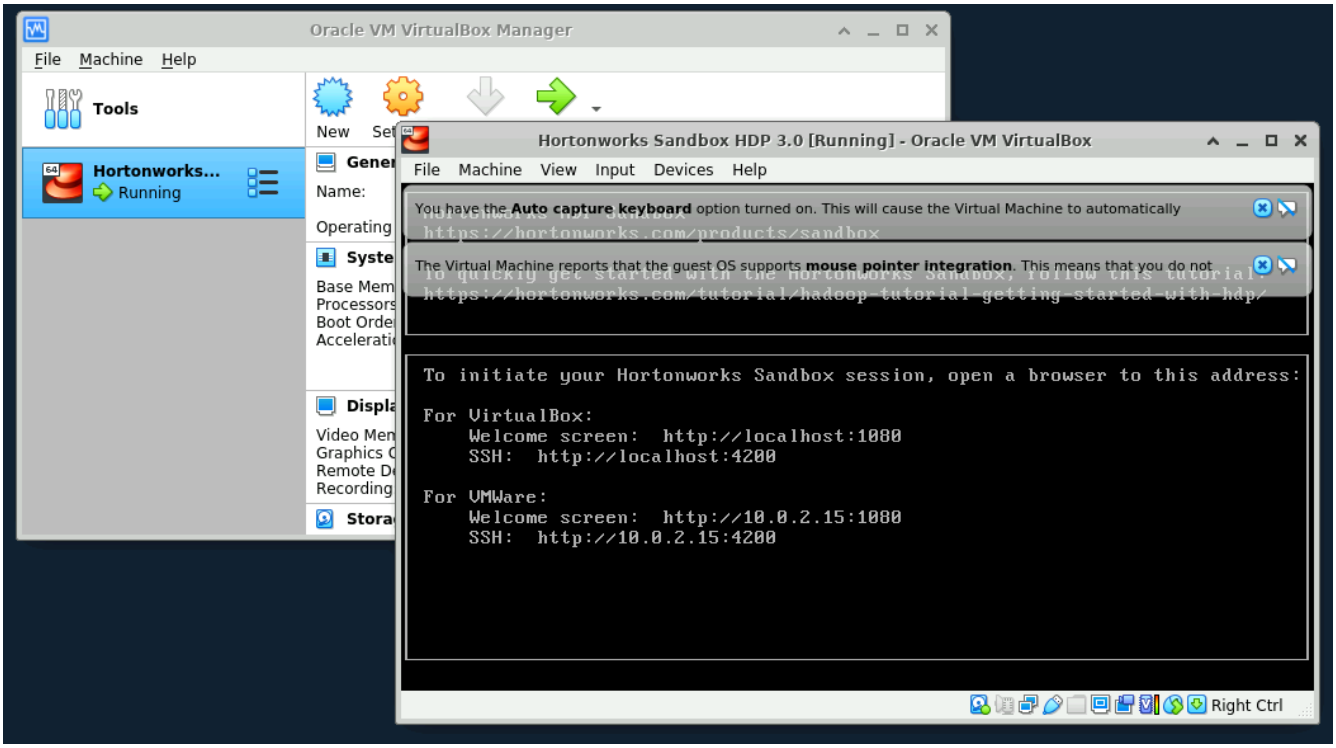

8. You are going to access the HPD VM via the browser *on the Nuvelab VM Desktop.*. Open the Chrome browser and go to the “localhost:1080.” This will bring up the HDP splash screen.

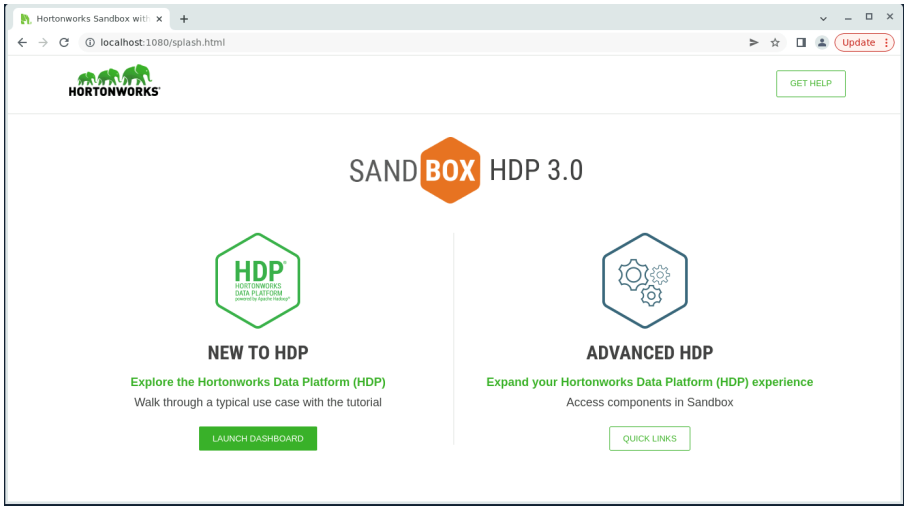

9. Click the option of the left. This will bring you to the Ambari management console. When prompted for a login use the credentials

   User name: **raj_ops** password: **raj_op**s

10. What you will see is the the display of the various services in the HPD starting up. It may take some time for them all to start.

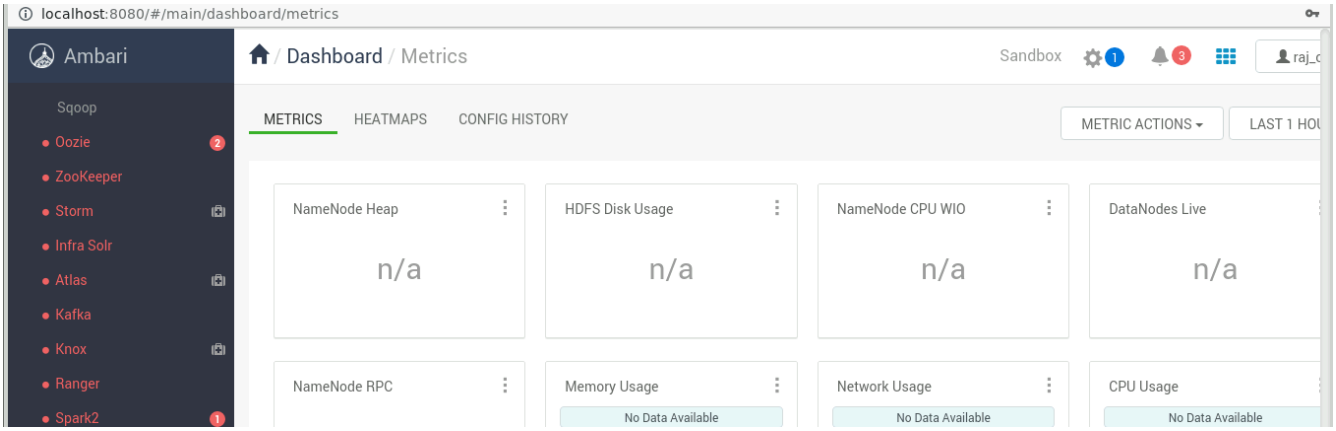

11. What you are waiting for is the job “start all services” to finish

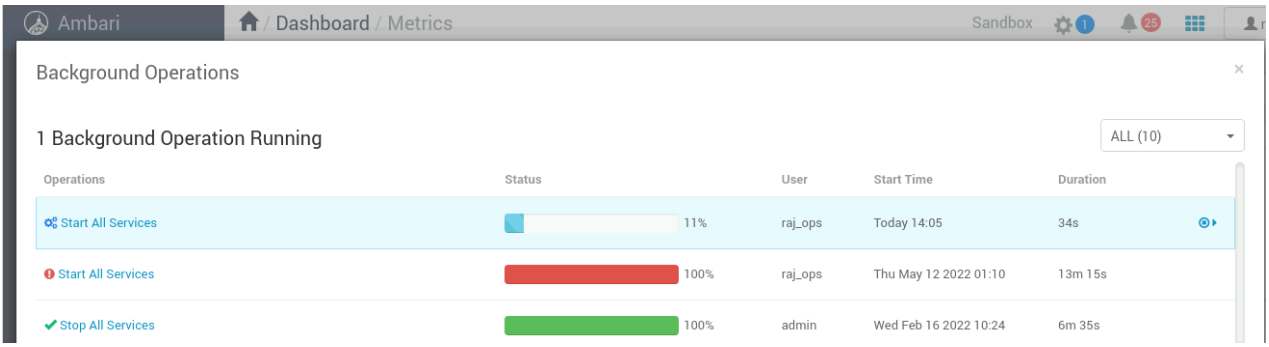

12. Once finished, there should be no red “alarm” icon.

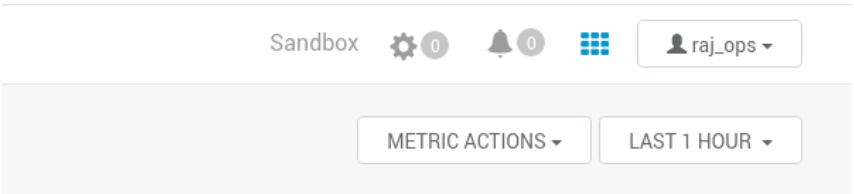

13. If there are any red alerts, you may need to use the drop down on the left hand side to restart all the services. If this doesn't work, then restart the HPD VM.

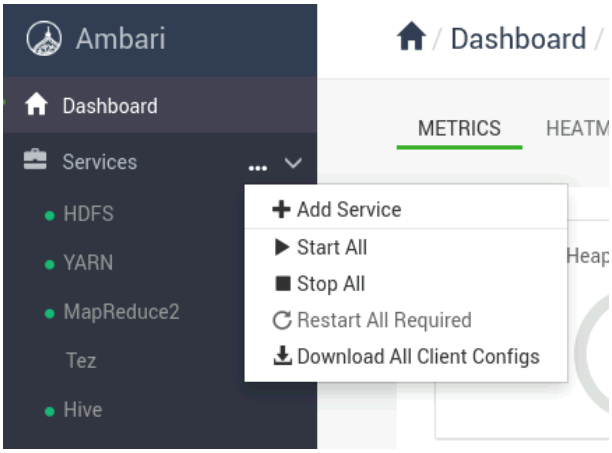

14. At this point, your HPD VM should be ready to use.

Hive Users and Labs

The HDP VM has two predefined users: maria_dev and raj_ops. The password for each user is the same as their user ID. If you have to do any configuration of the HPD environment, like setting the security, then you should log in as raj_ops since that ID has the access privileges needed. However, for the actual lab work, you should be logged in as maria_dev.

You can create your own user account if you want, but it is not necessary for the labs.

The lab materials are located in a git hub repository here:

https://github.com/ExgnosisClasses/VisaSQL101-June5

You can clone this repo to three different places.

1. Your Host Machine. This will enable you to keep a permanent copy of the repo but it will be inaccessible from the Nuvelab VM and the HDP VM.
2. The Nuvelab VM. This the recommended place to clone the lab since the data in the repo can be easily uploaded into the HPD Hadoop file system through the graphical in-terace. This is the options that will be used in all the lab instructions that follow.
3. The HDP VM. If you want to use the command line HDFS commands to load and work with the data, you can also clone the repository to your HPD VM.

Installing the repo on the Nuvelab VM

1. For convenience, you can copy data between your Host Machine and the Nuvelab VM.
2. To copy to the Nuvelab VM, press ctl-alt-shift on only the left side of the keyboard. This will open up a window on the left side of the screen. Anything you put into that window will be inserted into the Nuvelab VM. Once entered, press ctl-alt-shift on the left side of the keyboard to remove the window.

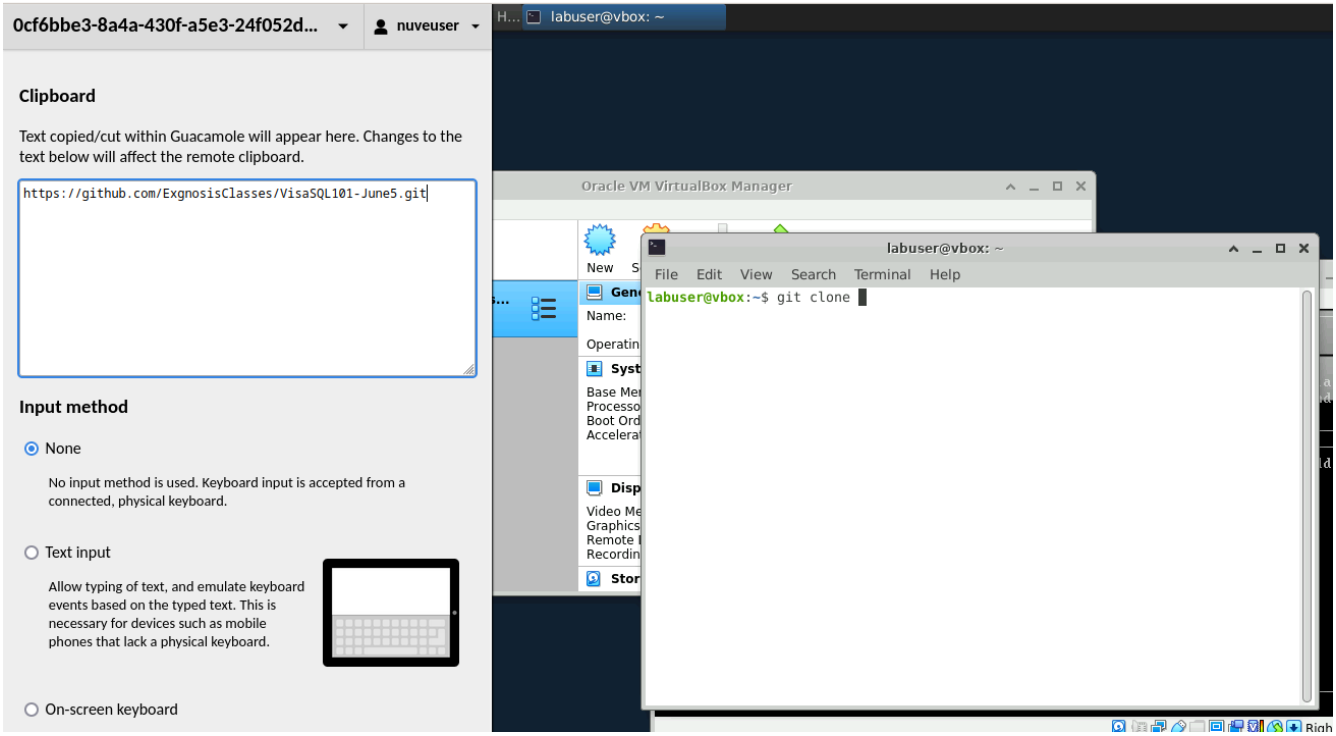

3. To add the repo to the Nuvelab VM, open up a terminal window from the menu at the bottom of the screen and execute the git clone command.

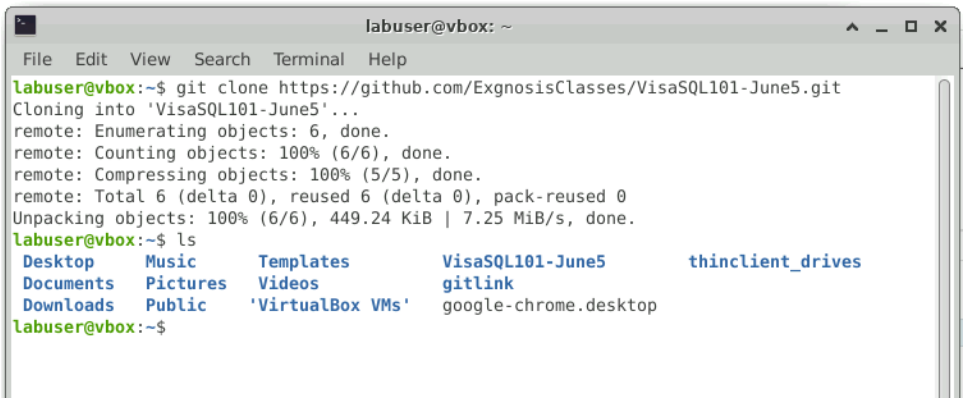

4.

Installing the repo in the HDP VM

1. First, open a web shell for the maria_dev user since that is the id that will be used in the labs in the HDP VM.
2. In a new browser window, go to localhost:4200. When prompted, enter maria_dev for both the user id and password.
3. Then you can execute the git clone command to install the repo in the HDP VM. *However, it is not necessary to do this to complete the labs.*

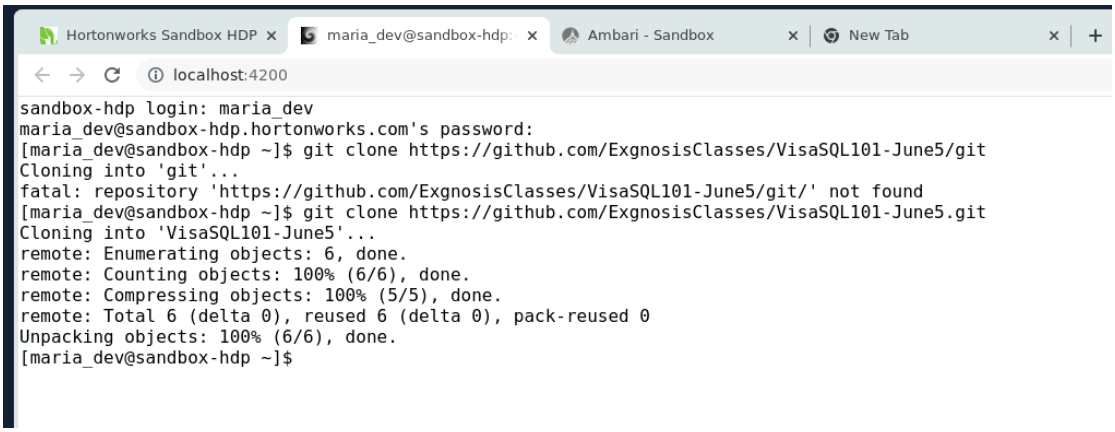

Starting Hive

1. First, open a web shell for the maria_dev user since that is the id that will be used in the labs in the HDP VM.
2. In a new browser window, go to localhost:4200. When prompted, enter maria_dev for both the user id and password.
3. At the command prompt, enter hive to start the hive shell.

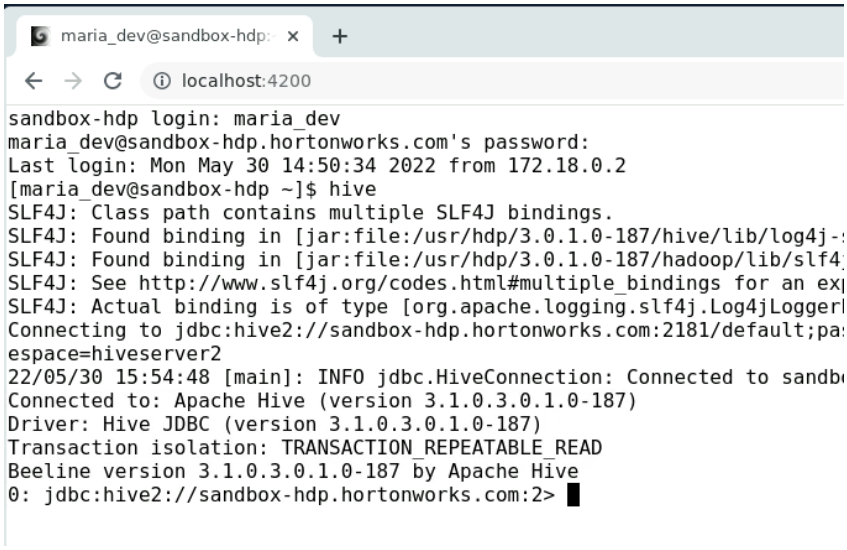

4. To ensure hive is running, execute the command “show databases;” which should return the output below.

```
0: jdbc:hive2://sandbox-hdp.hortonworks.com:2> show databases;
INFO  : Compiling command(queryId=hive_20220530155756_f62d66ad-5895-4756
INFO  : Semantic Analysis Completed (retrial = false)
INFO  : Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:dat
s:null)
INFO  : Completed compiling command(queryId=hive_20220530155756_f62d66ad
INFO  : Executing command(queryId=hive_20220530155756_f62d66ad-5895-4756
INFO  : Starting task [Stage-0:DDL] in serial mode
INFO  : Completed executing command(queryId=hive_20220530155756_f62d66ad
INFO  : OK
+----------------------+
|    database_name     |
+----------------------+
| default              |
| foodmart             |
| information_schema   |
| sys                  |
+----------------------+
4 rows selected (0.999 seconds)
```

Final Details

1. The README document on the Nuvelab VM desktop is incorrect. It gives the root password for the HDP VM as “labuser123$”. This is incorrect. The correct password is the same as for the Nuvelab VM “Nuvelabs123$”.
2. At the hive shell, run the following command.

   set hive.cli.print.header=true;

3. If you get an error like below, then you need to make the modifications described in the other pdf document “Security fix for labs.pdf”

```
0: jdbc:hive2://sandbox-hdp.hortonworks.com:2> set hive.cli.print.header=true;
Error: Error while processing statement: Cannot modify hive.cli.print.header at runtime. It is not in list of params that are allowed t
o be modified at runtime (state=42000,code=1)
0: jdbc:hive2://sandbox-hdp.hortonworks.com:2> set hive.cli.print.header;
+-----------------------------+
|             set             |
+-----------------------------+
| hive.cli.print.header=false  |
+-----------------------------+
1 row selected (0.159 seconds)
0: jdbc:hive2://sandbox-hdp.hortonworks.com:2>
```

When you are done, you should close the HDP VM window and select the option to power off the VM. Shown on the next page. This will ensure the next time you use the machine, all of the services will be initialized correctly. If you just suspend the machine, not all of the services may resume correctly leaving your VM in an unstable state.

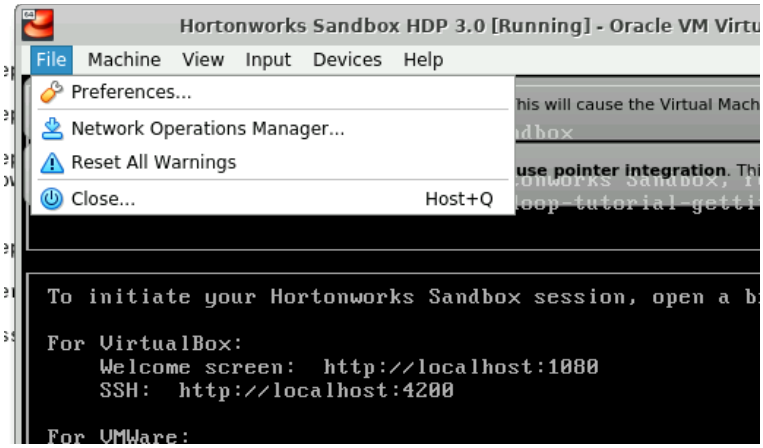

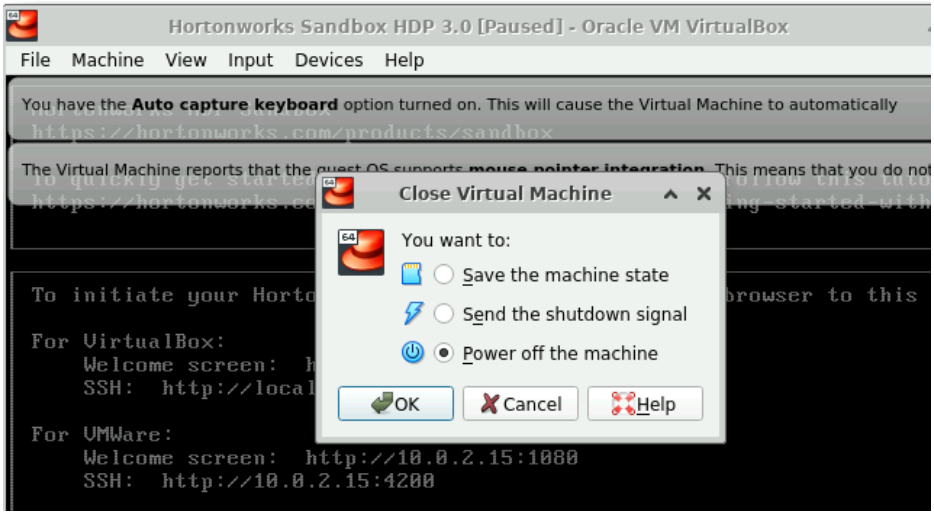
Hortonworks Sandbox HDP 3.0 [Paused] - Oracle VM VirtualBox
File Machine View Input Devices Help
You have the Auto capture keyboard option turned on. This will cause the Virtual Machine to automatically
https://hortonworks.com/products/sandbox
Close Virtual Machine
You want to:
Save the machine state
Send the shutdown signal
Power off the machine
OK
Cancel
Help
To initiate your Horto
browser to this
For VirtualBox:
Welcome screen:
SSH: http://local
For VMWare:
Welcome screen: http://10.0.2.15:1080
SSH: http://10.0.2.15:4200